Using Cell Phone Keyboards is ($\mathcal{NP}$) Hard

Peter Boothe

Manhattan College
`peter.boothe@manhattan.edu`

Abstract. Sending text messages on cell phones which only contain the keys 0-9 and # and * can be a painful experience. We consider the problem of designing an optimal mapping of numbers to sets of letters to act as an alternative to the standard $\{2 \rightarrow \{abc\}, 3 \rightarrow \{def\} \ldots\}$. Our overall goal is to minimize the number of buttons that must be pressed to enter an aferage message in English. We prove that the problem in almost all of its variations is $\mathcal{NP}$ hard, but describe several mappings which improve the standard one. With luck, one of these new mappings will achieve success similar to that of the Dvorak layout for computer keyboards.

Typing on a keyboard which has fewer keys than there are letters in the alphabet can be a painful task. There are a plethora of input schemes which attempt to make this task easier, but the one thing they all have in common is that all of these input methods use the standard mapping of numeric keys to alphabetic numbers of $\{2 \rightarrow \{abc\}, 3 \rightarrow \{def\}, 4 \rightarrow \{ghi\}, 5 \rightarrow \{jkl\}, 6 \rightarrow \{mno\}, 7 \rightarrow \{pqrs\}, 8 \rightarrow \{tuv\}, 9 \rightarrow \{wxyz\}\}$. In this paper, we consider schemes of rearranging the numbers on the keys to make messages easier to type. Most variants of this problem turn out to be NP-hard, unfortunately.

Before we get any farther, let us sketch the basic problem that we will keep revising and revisiting throughout this paper.

Problem 1 (MINIMUMKEYSTROKES).
INSTANCE *A set of letters corresponding to an alphabet A ($|A| = n$), a number of keys k, an input method IN, and a set of tuples of words and frequencies W. The frequencies in W are integers, and the words are made up of solely of elements of A. We will treat IN as a function which, given a partition of A and a word w, returns how many keystrokes are required to type w.*

QUESTION *What is the best partition of A into k sets, such that the total number of keystrokes to type every word in W its associated frequency times is minimized? Equivalently, what is the partition of P of A ($|P| = k$) that will minimize*

$$\sum_{(w,f) \in W} f * \mathrm{IN}(w, P)$$

Fig. 1. A cell phone keyboard

Over the course of this paper we consider three different real-world schemes for IN (basic typing, T9, and predictive T9), and for each variant that is proven $\mathcal{NP}$-hard, we consider the restriction on P that requires that we keep the alphabet in alphabetical order. In all cases where it is computationally feasible we provide results for the case of the English language, on 8-key keyboards, using the British National Corpus[1] (BNC). One feature of note is that no matter what scheme we use, the problem is trivial if the number of keys is not smaller than the number of letters in the alphabet ($k \geq |A|$). Using tiny keyboards is only (computationally) difficult when the keyboards don't have enough keys.

1   Setting a Baseline with the Easiest Problem

The baseline we compare against is the most painful method of text entry. In this method, to type an 'a', the user of the cellphone types the '2' key; to type a 'b' they type '22', to type 'c' they type '222', to type a 'd' they type '3', to type an 'e' they type '33', and so on. To type a word with multiple letters that use the same key, such as 'accept', the user must pause between keystrokes. Thus, the full key entry sequence for 'accept', using '.' to indicate a pause, is '2.22.223378'.

Initially, we will completely neglect the pauses and concern ourselves solely with the number of keystrokes required. This implies that 'a' will always require one button press, 'b' will always require two, and so on. Thus, to get a baseline of how many keystrokes are required to enter the entirety of our corpus of words W, we count the total number of occurences of each letter, and multiply that number

of occurrences by the number of button pushes required by that letter. Running this on the British National Corpus using the standard cell phone keyboard layout, we find that entering the entirety of the 100,106,029 word occurrences in the corpus would require 948,049,046 button presses.

If we examine a cell phone keyboard (Fig. 1) then we can see that there are some terrible choices being made. The frequently-occurring letters 'r' and 's' require more keystrokes than 'q'! Surely, a better designed keyboard can do better than this. If we just reversed the order of the letters on the 7 key, from 'pqrs' to 'srqp', the number of button presses required would drop to 865,118,331 — a savings of more than 8.7%! If we assume that every button press takes an equal amount of time, then this corresponds the users spending 8.7% less time entering their text messages. In an effort to find out how low this can go, we define the problem BASICCELLPHONETYPING based on the template problem on page 1.

Problem 2 (BASICCELLPHONETYPING).
INSTANCE   *A set of letters corresponding to an alphabet A ($|A| = n$), a number of keys k, and a set of tuples of words and frequencies W. The frequencies in W are integers, and the words are made up of solely of elements of A.*

QUESTION   *What is the best partition P of A into k sequences, such that the total number of button presses to type every word in W its associated frequency times is minimized? The number of button presses is equal to the order of the letter in the sequence assigned to a given key. For example, in the key $2 \to \{abc\}$, a requires one keystroke and c requires 3. Equivalently, if we denote the number of button presses required to type letter a with partition P as $\mathrm{IN}_P(a)$, then we want to find the P that minimizes*

$$\sum_{(w,f)\in W} \sum_{c\in w} \mathrm{IN}_P(c) * f$$

To solve BASICCELLPHONETYPING we construct a provably optimal greedy algorithm which works in time $O(|W|+|A|\log|A|)$. Our algorithm is detailed in Fig. 2, and involves creating a histogram of the letters, and then assigning letters to keys round-robin style in the order from most-popular to least-popular. To prove optimality, we invoke an exchange argument.

Theorem 3. *The greedy algorithm finds an optimal solution to* BASICCELL-PHONETYPING.

Proof. We will assume, for simplicity of proof, that all letters occur a different number of times in the corpus. We begin by comparing two assignments of letters to keys by, for each assignment, constructing a sequence of sets, or spectrum, S. The set S_1 (layer 1) is the set of all letters which require a single button press (that is, they are the first letter in their sequence on their assigned key), S_2 (layer 2) is the set of all letters which require two button presses, and so on. If two

GreedyBasicCellPhoneTyping (A, k, W)
    // A is the set of letters
    // k is the number of keys
    // W is a set of pairs of words and their corresponding frequencies
    lettercount $\leftarrow$ new_map()
    for $c \in A$
        lettercount$[c] \leftarrow 0$
    for (word, frequency) $\in W$
        for $c \in$ word
            lettercount$[c] \leftarrow$ lettercount$[c]$+frequency
    ordered $\leftarrow$ [(lettercount$[c], c$) **for** $c \in$ lettercount]
    sort(ordered)
    keys $\leftarrow$ a array of length k where each element is an empty list
    for $i \leftarrow 0 \ldots length$(ordered)
        (count, char) $\leftarrow$ ordered$[i]$
        append(char, keys$[i mod k]$)
    return keys

Fig. 2. The greedy algorithm for Prob. 2.

assignments have equal spectra, then one assignment may be transformed into the other assignment simply by swapping letter between keys without increasing or decreasing the total number of button pushes required to enter the corpus of words. Two assignments are isomorphic iff their spectra are equal.

Now assume for the sake of contradiction that we have assignment D, with spectrum T, that is not isomorphic to the greedy solution G with spectrum S. Note that in spectrum resulting from the greedy algorithm, all letters at layer i occur strictly more frequently than the letters at layer j, $j > i$. Also, in S, for all layers i except for the last layer, $|S_i| = k$. Because T and S are not isomorphic, in T there must be at least one layer i such that $T_i \neq S_i$. Let T_i be the first layer for which that is true.

There are two cases for layer T_i. In the first case, $|T_i| < k$, and we may remove any element from a later layer place it in layer T_i and create a strictly better layout. In the second case, there is some letter a such that $a \in T_i$ and $a \notin S_i$. We choose the least-frequent such a. Because the greedy layout algorithm creates layers in order of frequency, we know that the frequency of a is less than that of some letter b in layer T_j, $j > i$. Therefore, by creating a new layout with a and b swapped between T_i and T_j, we have strictly decreased the number of button presses. Therefore, in both cases, T was not an optimal layout, and we have reached our contradiction. $\qquad\square$

For the BNC, the optimal layout has the spectrum

$$[\{etaoinsr\}, \{hldcumfp\}, \{gwybvkxj\}, \{qz\}]$$

and any layout with that spectrum requires 638,276,114 button presses to entire the entire BNC instead of our original requirements of 948,049,046. This represents a savings of 32.67% over our original layout, but only for the users who

type using this most basic of input methods. Unfortunately, this is the group of users who are likely the least adaptable to change, as almost all "digital natives" use predictive methods to input text messages[1]. In order to place the least burden of "newness" on these users whilst still decreasing the number of button presses, we now consider schemes where the only alteration of the keyboard is the rearrangement of the sequence of letters for a given key.

By an argument symmetric to the proof of Theorem 3, we find that the optimal layout in this new scheme is to place the letters of a given key in sorted order, according to frequency. Thus, the keyboard layout changes to $2 \rightarrow [acb], 3 \rightarrow [edf], 4 \rightarrow [ihg], 5 \rightarrow [lkj], 6 \rightarrow [onm], 7 \rightarrow [srpq], 8 \rightarrow [tuv], 9 \rightarrow [wyxz]$ and requires 678,547,463 button presses to enter the whole corpus. This represents a 28.43% savings, but has the added benefit that it does not change which key is mapped to which set of letters. This makes it **legacy preserving**, as creating a keyboard with this layout will not invalidate such advertising gems as 1-800-FLOWERS. Thus, we can speed up users by 28.43% (neglecting interletter pauses) and not undermine 50 years of advertising. This layout represents perhaps the most plausible layout yet[2]. After setting up this baseline for the easy problem, we turn our attention to optimizing cellphone keyboards for predictive input methods.

2   The T9 Input Method

In an effort to minimize the pain of cellphone keyboard typing, cellular telephone manufacturers have created the T9 input method, which attempts to guess which letter (of the 3 or 4 possible) is intended when a user hits a single key. Enhancements of this input method also provide **speculative lookahead** to report, at any given time, what word it is that the user is most likely trying to enter and provide a completion. Unfortunately, because there are fewer keys on the cellphone keyboard than there are letters in the English alphabet, there are words which have different spellings but the same input sequence. As an example, in the traditional mapping of keys to characters both 'me' and 'of' have the input sequence '63'. Two words with the same input sequence force the user to press a third button to cycle through the possibilities in order from most likely ('of') to least likely ('me'). Even worse: 'home', 'good', 'gone', 'hood', 'hoof', 'hone', and 'goof' all have the input sequence '4663'. If the * key is used as the "next match" button, then the user will actually have to type in '4663******' to type in the word 'goof', for a total of 10 button presses — exactly the same number of button presses required using the default layout and the basic input method, and

---

[1] I have no statistics on this except for an informal survey of one of my classes. In that class, every student who used SMS either had a cellphone with a 26+ key alphabetic keyboard or used a predictive method.

[2] The described layout is the only legacy preserving layout in this paper, and therefore should be considered the most practical suggestion. It also has the added benefit of not angering any of the organizations behind the numbers 1-800-FLOWERS, 1-800-THRIFTY, and 1-800-MARINES

three more button presses than is required when using the basic input method with an optimized layout.

When two words have the same input sequence (neglecting the *'s at the end) then these words are **t9onyms**. When two or more words are t9onyms, then the less-popular words require extra keypresses, raising the expected number of keypresses to type in our corpus. In order to type a given word, one must press one button for every character in the word, followed by pressing the * key as many times as there are t9onyms which are more likely than the desired wored. Because typing on a cellphone keyboard is already an unpleasant experience, we would like to minimize the expected number of keystrokes. The number of keystrokes a word requires is equal to the number of letters in the word, plus the number of t9onyms which are more frequent than the word. Formally, we extend Prob. 1 and define the MINIMUMT9KEYSTROKES problem as:

Problem 4 (MINIMUMT9KEYSTROKES).
 INSTANCE *An alphabet A, a set of words and their associated frequencies W, and a number of keys k ($|A| > k$).*

QUESTION *What is the partition P of A into k sets which minimizes the total button presses required to enter the entire corpus using the T9 input method?*

and the corresponding decision problem is

Problem 5.
 INSTANCE *A set A, a set W of sequences of elements of A, a mapping f from sequences in W to the integers, and a number t.*

QUESTION *Is there a partition P of A into k sets such that, if we denote $P(w)$ as the sequence of partitions $[p_i|w_i \in p_i, p_i \in P]$,*

$$\sum_{w \in W} (\text{len}(w) + \text{order}(w, P, W, f)) * f(w) \leq t$$

where len(w) is the length of the sequence w and order(w, P, W, f) is the size of $\{s \in W | P(s) = P(w) \wedge f(s) > f(w)\}$, which is the set of sequences which map to the same sequence of partitions, but are strictly more popular.

Note that this problem is in $\mathcal{NP}$, as any assignment may be verified to require fewer than f button pushes in time proportional to the total length of all the words in W. In order to prove completeness, however, we first prove the $\mathcal{NP}$-completeness of an intermediate problem: UNIQUEPATHCOLORING.

Problem 6 (UNIQUEPATHCOLORING).
 INSTANCE *A graph $G = (V, E)$, a set of paths P, and a parameter k. A path p is a sequence of vertices in which adjacent vertices in p are also adjacent in the edge set E.*

QUESTION *Is there a k-coloring of G such that every path $p \in P$ has a unique*

coloring? If we consider the coloring function $\chi(v)$ to map vertices to colors, then we can extend this notation by having $\chi(p)$ map a path $p = [v_1, v_2, \ldots]$ to the sequence $\chi(p) = [\chi(v_1), \chi(v_2), \ldots]$. Is there a coloring χ of V such that

$$|\{\chi(p) \; \forall p \in P\}| = |P| \; ?$$

Theorem 7. UNIQUEPATHCOLORING *is $\mathcal{NP}$-complete.*

Proof. To prove that UNIQUEPATHCOLORING is $\mathcal{NP}$-complete, we begin by noting that any coloring of V may be verified, in polynomial time, to map each path to a unique sequence of colors. Therefore, the problem is in $\mathcal{NP}$. To prove completeness, we reduce from GRAPHCOLORING[**?**].

An instance of GRAPHCOLORING consists of a graph $G = (V, E)$ and a parameter k. We then ask the question of whether there is a k-coloring χ of the vertices of the graph such that $\forall (u, v) \in E$, $\chi(u) \neq \chi(v)$. We transform an instance of GRAPHCOLORING into an instance of UNIQUEPATHCOLORING in the following manner:

Given $G = (V, E)$ and k from GRAPHCOLORING, we create

$$G' = (V \cup \{0, 1\}, E \cup \{(0, 1)\} \cup \{(v, 0) \; \forall v \in V\})$$

We then uniquely number each edge in E with the numbers $1 \ldots |E|$. For each edge $e = (u, v)$ numbered i, we create the path set

$$p_i = \{[v, 0, b_1(i), b_2(i), b_3(i), \ldots b_{\lceil \log_2 i \rceil}],$$
$$[u, 0, b_1(i), b_2(i), b_3(i), \ldots b_{\lceil \log_2 i \rceil}]\}$$

where $b_1(i)$ is the first digit of i in binary, $b_2(i)$ is the second digit of i in binary, and so on. Now we create our path set

$$P = \{[0, 1], [1, 0]\} \cup \bigcup_{i=1}^{|E|} p_i \; .$$

We then ask the question of whether G' and P can be unique-path colored using only k colors.

If there is a k-coloring χ of G, then we use that coloring to generate a k-unique-path coloring of G' and P in the following manner: First, assign vertices 0 and 1 different colors from the range of χ. Next, assign each vertex to the color it received in the k-coloring of G'. Now, every element of our path set corresponds to a unique color sequence. $[0, 1]$ and $[1, 0]$ are the only sequences of length two, and because we assigned these two vertices different colors, $\chi([0, 1]) \neq \chi([1, 0])$. If path x and path y are not from the same p_i, then they have a different sequence of zeroes and ones. Therefore, because 0 and 1 are assigned different colors, the only possible path that a given path in some p_i might be confused with is the other path in that p_i. But each p_i corresponds to a an edge in E, and we know for all edges in $(u, v) \in E$ that $\chi(u) \neq \chi(v)$. Therefore, both paths within a

given p_i also have a distinct color sequence. Therefore, given a k-coloring of G, we can generate a k-unique-path coloring of G' and P.

Now we prove that given a k-unique-path coloring of G' and P we can create a k-coloring of G. Given a k-unique path coloring, we are guaranteed that no path in P has the exact same coloring as any other path in P. This means that, for all i, the two paths in p_i are distinct. The only difference between the two paths in p_i is in the first vertex of the path. Therefore, for every p_i, the vertices of corresponding edge in must be given distinct colors, and this is true for all i and $e \in E$. Therefore, for all $(u, v) \in E$, the color assigned to u is distinct from the color assigned to v. Therefore, from a k-unique-path coloring of G' and P, we have a k-coloring of $G = (V, E)$.

□

Now that we have proven UNIQUEPATHCOLORING to be $\mathcal{NP}$-complete, we immediately use it in a reduction.

Theorem 8. MINIMUMT9KEYSTROKES *(the decision problem form, Problem 5) is $\mathcal{NP}$-complete.*

However, our arbitrary remappings of the keyboard may be quickly deemed unusable. Therefore, we define a refinement of the problem, in which the characters are required to be kept in alphabetical order as a sequence, and our only choices are about how to divide the subsequence into key assignments. The default layout can be described as the partition $abc|def|ghi|jkl|mno|pqrs|tuv|wxyz$. Is this partition optimal? In order to decide this, we define the problem MINIMUMKEYSTROKEPARTITION as:

Problem 9 (MINIMUMKEYSTROKEPARTITION**).**
 INSTANCE   *A sequence of letters $A = \{a_1, a_2, \ldots\}$, a set of words and their associated probabilities W, and a set of keys $K = \{2, 3, 4, \ldots\}$ ($|A| > |K|$).*

QUESTION   *What is the mapping f of $A \to K$ which minimizes the expected number of characters to type a word from W, and where $f(a_1) = 2$ and if $f(a_i) = k$, then either $f(a_{i+1}) = k$ or $f(a_{i+1}) = k + 1$?*

As before, the corresponding decision problem is

Problem 10.
INSTANCE   *A sequence of letters $A = \{a_1, a_2, \ldots\}$, a set of words and their associated probabilities W, a set of keys $K = \{2, 3, 4, \ldots\}$ ($|A| > |K|$), and a parameter e.*

QUESTION   *Is there a mapping f of $A \to K$ in which the expected number of characters to type a word from W is less than e, and where $f(a_1) = 2$ and if $f(a_i) = k$, then $f(a_{i+1}) \in \{k, k+1\}$?*

Again, we note that this decision problem is in $\mathcal{NP}$, as any proposed partition may be verified, in polynomial time, to have an expected value less than e.

3   Exhaustive Search

Because our problems are $\mathcal{NP}$-hard, we move on to combinatorial enumeration and experimental results. The number of partitions of a sequence of size n into k subsequences is equal to $\binom{n-1}{k-1}$, and in the particular case of the 26 letter alphabet and the eight keys available on a cell phone keyboard, we find that there are $\binom{26-1}{8-1} = 480,700$ possible sequences.

References

1. British National Corpus Consortium. British National Corpus, 2000.